

WEEK 5: TEXT MINING 3

SECU0057

BENNETT KLEINBERG

13 FEB 2020

Applied Data Science

WEEK 5: TEXT MINING 3

TODAY

- text similarity
- intro to word embeddings

What makes two strings similar?

FACTORS IN SIMILARITY

- lexical similarity
- phonetic similarity
- semantic similarity

Underlying challenge: **quantification**

STRING SIMILARITY

- Levenshtein distance
- Damerau-Levenshtein distance
- q-gram distances

Aim: numerical measure of distance

```
string_a = 'crime'  
string_b = 'criminal'
```

How much do these two differ?

LEVENSHTEIN'S APPROACH

Distance between A and B :

1. change A to obtain B
2. allowed edits:
 - substituting characters
 - adding characters

LEV('CRIME', 'CRIMINAL')

1. crim_ei
2. crimin
3. crimina
4. criminal

4 edits are required to change 'crime' into 'criminal'.

IN R

```
library(stringdist)

stringdist(a = 'crime'
           , b = 'criminal'
           , method = 'lv')
```

```
## [1] 4
```

$$Lev_{(crime,criminal)} = 4$$

YOUR TURN

Pair 1:

- A = “london”
- B = “amsterdam”

Pair 2:

- C = “london”
- D = “condom”

```
stringdist(a = 'london'  
           , b = 'amsterdam'  
           , method = 'lv')
```

```
## [1] 8
```

```
stringdist(a = 'london'  
           , b = 'condom'  
           , method = 'lv')
```

```
## [1] 2
```

CONVERTING DISTANCES TO SIMILARITY

1. calculate distance $dist$ (e.g. $Lev_{(crime,criminal)} = 4$)
2. calculate maximum distance $\max dist$
3. obtain quotient of $\frac{dist}{\max dist}$
4. subtract quotion from 1: $sim = 1 - \frac{dist}{\max dist}$

DISTANCE TO SIMILARITY

1. $Lev_{(crime,criminal)} = 4$

2. $\max Lev_{(crime,criminal)} = 8$

3. $\frac{Lev_{(crime,criminal)}}{\max Lev_{(crime,criminal)}} = \frac{4}{8}$

4. $1 - \frac{Lev_{(crime,criminal)}}{\max Lev_{(crime,criminal)}}$

$$1 - \frac{4}{8} = 0.5$$

USING STRINGSIM(. . .)

```
stringsim(a = 'crime'  
          , b = 'criminal'  
          , method = 'lv')
```

```
## [1] 0.5
```

Q-GRAMS

- essentially: n-grams on character-level

```
qgrams('crime', 'criminal', q = 2)
```

```
##      cr ri im me mi na in al
## V1  1  1  1  1  0  0  0  0
## V2  1  1  1  0  1  1  1  1
```

SETTING Q

Note that q is analogous to n in n-grams:

```
qgrams('crime', 'criminal', q = 3)
```

```
##      cri rim ime imi min nal ina
## v1    1   1   1   0   0   0   0
## v2    1   1   0   1   1   1   1
```

Q-GRAMS FOR STRING DISTANCE

Simple approach: absolute difference

```
stringdist(a = 'crime'  
           , b = 'criminal'  
           , method = 'qgram'  
           , q = 3)
```


```
## [1] 5
```

DISTANCES BETWEEN VECTORS

Suppose we have got two vectors:

- $\vec{v}_1 = [1, 2]$
- $\vec{v}_2 = [4, 3]$

Vectors $[1,2]$ and $[4,3]$




EUCLIDEAN DISTANCE

Uses Pythagorean theorem.

For two two-dimensional locations:

- build a right triangle
- use $a^2 = b^2 + c^2$ to calculate the length of the hypotenuse c

Vectors $[1,2]$ and $[4,3]$



BY HAND:

- $a = x_2 - x_1 = 4 - 1 = 3$
- $b = y_2 - y_1 = 3 - 2 = 1$
- $c^2 = a^2 + b^2$

Thus:

$$c = \sqrt{a^2 + b^2} = \sqrt{9 + 1} = 3.16$$

EUCLIDEAN DISTANCE IN 3 DIMENSIONS

$$\textit{dist}(A, B) = \sqrt{(A_1 - B_1)^2 + (A_2 - B_2)^2 + (A_3 - B_3)^2}$$

ON THE Q-GRAMS

```
qgrams('london', 'condom', q = 2)
```

```
##      lo on nd om do co
## V1   1  2  1  0  1  0
## V2   0  1  1  1  1  1
```

VECTORS V1 AND V2

```
matrix_1 = matrix(data = c(c(1,2,1,0,1,0),c(0,1,1,1,1,1))  
                    , nrow = 2, byrow = T)  
  
dist(matrix_1, method = 'euclidean')
```


```
##      1  
## 2 2
```

EUCLIDEAN DISTANCE AND MAGNITUDE

- takes into account the magnitude of vectors
- but this is not always meaningful
- different metric: *cosine* distance

COSINE DISTANCE

Cosine distance: about the angles



COSINE SIMILARITY

$$csim = \frac{A \times B}{\sqrt{A \times A} * \sqrt{B \times B}}$$

_Note: $A \times B$ is the dot product of the vector.

IN R

```
V1 = c(4,2,3)
V2 = c(1,3,1)

cossim = function(A, B){
  numerator = sum(A*B)
  denominator = sqrt(sum(A*A))*sqrt(sum(B*B))
  cosine_sim = numerator/denominator
  return(cosine_sim)
}

cossim(V1, V2)
```

```
## [1] 0.7278603
```


COSINE SIMILARITY ON Q-GRAMS

```
qgrams('london', 'condom', q = 2)
```

```
##      lo on nd om do co  
## v1  1  2  1  0  1  0  
## v2  0  1  1  1  1  1
```

```
stringsim('london', 'condom', q = 2, method = 'cosine')
```

```
## [1] 0.6761234
```

SIMILARITY OF PHRASES

```
phrase_1 = "I think numbers are great"  
phrase_2 = "Numbers are useless"  
  
stringsim(phrase_1, phrase_2, q=3, method = 'cosine')
```

```
## [1] 0.4551496
```

PROBLEMS WITH LEXICAL SIMILARITY?

CONSIDER THESE

```
[sand, beach]  
[water, sea]  
[bank, money]  
[euro, dollar]
```

WHAT ABOUT THE SEMANTIC DIMENSION?

- lexical dimension doable
- But what about the meaning of words?
- “sand” and “beach” are close to one another
- ... but lexical metrics fail to uncover this

SEMANTIC SIMILARITY

If a method can do this, then we expect:

- $\text{cossim}(\text{sand}, \text{beach}) > \text{cossim}(\text{bleach}, \text{beach})$
- $\text{cossim}(\text{euro}, \text{dollar}) > \text{cossim}(\text{neuro}, \text{euro})$

WORD EMBEDDINGS

- meaning of words determined by information associated with it
- computationally: we express each word as a vector
- that vector contains the information associated with the word

Does that work? When?

VECTORS ARE EVERYWHERE

“these are word embeddings”

```
## Warning: 'as.data.frame.dfm' is deprecated.  
## Use 'convert(x, to "data.frame")' instead.  
## See help("Deprecated")
```

document	these	are	word	embeddings
text1	1	0	0	0
text2	0	1	0	0
text3	0	0	1	0
text4	0	0	0	1

Every word is a vector (called: a **one-hot vector**)

EVERY DOCUMENT IS A VECTOR

```
## Warning: 'as.data.frame.dfm' is deprecated.  
## Use 'convert(x, to "data.frame")' instead.  
## See help("Deprecated")
```

document	career	upon	which	i	am	about	to	enter	appear	f
1825-Adams	2	16	27	15	4	1	101	1	1	
1829-Jackson	0	0	12	12	0	1	53	0	1	
1833-Jackson	0	5	14	6	2	1	46	0	0	

VECTOR-BASED WORD SIMILARITY

So can't we just calculate the vector similarity then?
(e.g. cosine similarity)

Ideas?

PROBLEM OF SPARSITY

Remember: dot products!

Dot product of two one-hot vector is always zero.

EMBEDDING A VECTOR

- we want a denser representation of each word
- two approaches:
 - word2vec: $P(\text{word}|\text{context})$ and $P(\text{context}|\text{word})$
 - glove: corpus co-occurrences
 - Details available [here](#)

EMBEDDINGS: AIM

From high-dimensional vector space to lower dimensional space.

E.g. from a one-hot vector of size 10,000 to an embedding of size 300.

PROPERTIES OF WORD EMBEDDINGS

- we have an embedding of each word
- that captures some kind of “context”
- each embedding in a given model has the same length

USING WORD EMBEDDINGS

- no need to build the embeddings
- typically we use pre-trained models (e.g. from Google, Facebook)
- e.g. Word vectors for 157 languages and the Glove embeddings
- we can then retrieve the embeddings for a given word

IN R

Note: heavy on your RAM (loads massive file into memory)

1. initialising pre-trained models
2. calculating similarity between terms

Code: https://github.com/ben-aaron188/r_helper_functions/blob/master/init_glove.R

OUR EXPECTATIONS (1)

cossim(sand, beach) > cossim(bleach, beach)

```
sand_emb = as.vector(glove.pt[row.names(glove.pt) == 'sand', ])  
beach_emb = as.vector(glove.pt[row.names(glove.pt) == 'beach', ])  
bleach_emb = as.vector(glove.pt[row.names(glove.pt) == 'bleach', ])
```


OUR EXPECTATIONS (2)

$$\text{cossim}(\text{euro}, \text{dollar}) > \text{cossim}(\text{neuro}, \text{euro})$$

```
euro_emb = as.vector(glove.pt[row.names(glove.pt) == 'euro', ])  
dollar_emb = as.vector(glove.pt[row.names(glove.pt) == 'dollar', ])  
neuro_emb = as.vector(glove.pt[row.names(glove.pt) == 'neuro', ])
```

```
cossim(euro_emb, dollar_emb)  
#[1] 0.7328042
```

```
cossim(euro_emb, neuro_emb)  
#[1] -0.08723018
```

ARITHMETICS WITH WORD EMBEDDINGS

What if we could do maths with meaning?

Do semantic relationships hold true in embeddings?

BERLIN - GERMANY + FRANCE

$$\text{If: } \overrightarrow{BERLIN} \approx \overrightarrow{GERMANY} - \overrightarrow{FRANCE}$$

$$\frac{\overrightarrow{GERMANY}}{\overrightarrow{BERLIN}} \therefore \frac{\overrightarrow{FRANCE}}{?}$$

```
berlin = as.vector(glove.pt[row.names(glove.pt) == 'berlin', 1])
germany = as.vector(glove.pt[row.names(glove.pt) == 'germany', 1])
france = as.vector(glove.pt[row.names(glove.pt) == 'france', 1])
```

```
mystery_1 = berlin - germany + france
```

```
[1]  0.5129700 -0.7331210 -0.1134250  0.4532580  0.5920800  0.5046900  
[7]  0.1716500  0.6980830  0.0461335  0.2364360 -0.2775940 -0.0869580  
[13] -0.2696400  0.1602300 -0.9382800  0.3996800  0.1305250 -0.6906400  
[19] -0.6324500 -0.3029900  0.2935000  0.1269000 -0.1562700 -1.1325400  
[25] ...
```

CLOSEST NEIGHBOURS TO THE MYSTERY VECTOR?

```
head(sort(cos_sim_vals[,1], decreasing = TRUE), 10)
```

mystery_1	paris	france	french	prohertrib	berlin	bruss
1.0000000	0.8827144	0.7558026	0.7075165	0.6943174	0.6665562	0.6574
lyon	london	le				
0.6526201	0.6407975	0.6403628				

KING - MAN + WOMAN

$$\frac{\vec{MAN}}{\vec{KING}} :: \frac{\vec{WOMAN}}{?}$$

```
head(sort(cos_sim_vals[,2], decreasing = TRUE), 10)
```

```
   mystery_2      king      queen  monarch  throne  daughter  prince  
1.0000000 0.8551837 0.7834414 0.6933802 0.6833110 0.6809082 0.6713142 0.6  
   mother elizabeth  
0.6579325 0.6563301
```

WHY DOES THIS WORK?

Possible explanation: the distributional hypothesis (Harris, 1954)

Words in that occur in the same context have the same meaning.

CAUTIONARY NOTE

Word embeddings are very powerful, but:

- instable
- numerically meaningless
- Limitations of word embeddings, Burdick (2019)

(see the required reading for today)

WHAT'S NEXT?

- Today's tutorial: text similarity, word embeddings
- Homework: more word embeddings, custom functions

Next week: reading week (then: Machine Learning 1)