


# neural networks

an introduction




**mite**





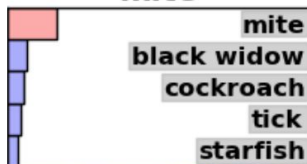
**container ship**



**motor scooter**



**leopard**



**grille**






**mushroom**



**cherry**



**Madagascar cat**



# background

2012 - Geoffrey Hinton (University of Toronto)

Achieved 37.5% error (Top 1), 17% error (Top 5) on ImageNet database.

60 million parameters, 650,000 neurons

From then on: increased computational power helped with the popularity of neural nets (GPUs)

# goal

- Motivate the structure
- What it means for a machine to learn
- What it means for learning to be 'deep'
- Introduce these ideas using a vanilla neural network

# approach

Classifier  
needs ↓


IT has to  
go through ↓

Then  
↓

Learning = Representation + Evaluation + Optimization


(Pedro Domingos, A Few Useful Things to Know about Machine Learning  
<https://homes.cs.washington.edu/~pedrod/papers/cacm12.pdf>)

# the inspiration




# not a neuroscientist, but:

- Each neuron receives inputs from other neurons
- The effect of each input line on the neuron is controlled by a synaptic weight
  - positive or negative
- The synaptic weights **adapt** so that the whole network learns to perform useful computations
  - recognising objects, language, movement, etc.
- We have about 86 billion neurons




# basic perceptron

Type of artificial neuron






# basic perceptron




# basic perceptron



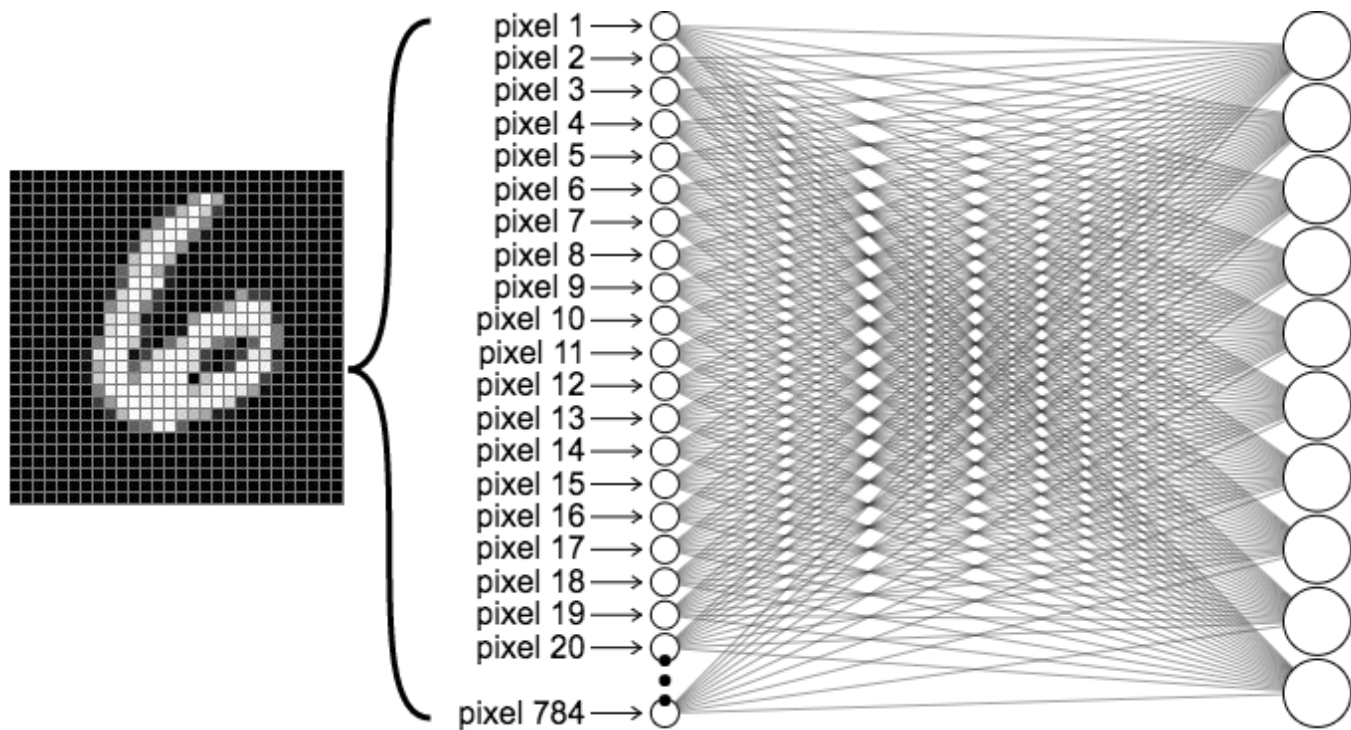
$$\text{output} = \begin{cases} 0 & \text{if } \sum_j w_j x_j \leq \text{threshold} \\ 1 & \text{if } \sum_j w_j x_j > \text{threshold} \end{cases}$$

# combining

deep neural network simply means there are multiple hidden layers!




idea




# basic neuron

in our simplified model, a neuron simply holds a number



**activation function: smooth**



*Sigmoid*

$$f(x) = \frac{1}{1 + e^{-x}}$$

# learning

finding the right weights

$$\sigma(w_1 a_1 + w_2 a_2 + w_3 a_3 \dots w_n a_n)$$

and biases

$$\sigma(w_1 a_1 + w_2 a_2 + w_3 a_3 \dots w_n a_n - 10)$$

bias

# what we have so far

- neurons
- layers of neurons with weights and biases
- activation function

what more do we need?

- some way to measure performance of our weights and biases
- some way to update our parameters based on how well they performed



# measuring performance

adjust weights and threshold to maximise performance.

intuitively,

$$\text{error} = \text{desired output} - \text{actual output}$$

# measuring performance

adjust weights and threshold to maximise performance.

intuitively,

$$\text{error} = \text{desired output} - \text{actual output}$$

$$\text{error} = || \text{desired output} - \text{actual output} ||$$

# measuring performance

adjust weights and threshold to maximise performance.

intuitively,

$$\text{error} = \text{desired output} - \text{actual output}$$

$$\text{error} = ||\text{desired output} - \text{actual output}||$$

$$\text{error} = (\text{desired output} - \text{actual output})^2$$


# measuring performance

Input	actual	Desired	Absolute Error	Square Error
0	0	0	0	0
1	3	2	1	1
2	6	4	2	4
3	9	6	3	9
4	12	8	4	16
Total:	-	-	<b>10</b>	<b>30</b>

$$\text{MSE} = \frac{1}{n} \sum_{i=1}^n (Y_i - \hat{Y}_i)^2.$$


# gradient descent

- now that we have the error, what we want to do is find the combination of weights and biases that minimise it





# gradient descent

- now that we have the error, what we want to do is find the combination of weights and biases that minimise it




# gradient descent

- now that we have the error, what we want to do is find the combination of weights and biases that minimise it




# gradient descent





# backpropagation

- gradient descent is part of a process called backpropagation
- first we calculate for each neuron how much we want to change, and then send this backwards



# what do we have now

recall: Learning = Representation + Evaluation + Optimization

- representation: neural network
- evaluation: mean squared error
- optimization: gradient descent + backpropagation

**tensorflow playground**

# neural networks in R

'neuralnet' package